Three-Tree Genetic Programming for Skill-Based Peer Coordination in Multi-Site Inventory Control

Author Name

***Abstract*—Multi-site inventory control requires local replenishment decisions and coordinated lateral transfers under uncertain demand. Treating every transfer pair independently can allocate the same donor stock more than once, while giving a central policy unrestricted access to all site states obscures the information requirements of coordination. This paper develops a reproducible method in which each site is a peer inventory agent with local observations, explicit decision skills, and a reservation ledger. An offline genetic programming (GP) search evolves a shared individual with exactly three expression trees: replenishment, support request, and support offer. Fixed skill wrappers interpret these expressions as bounded quantities. A non-evolved protocol exchanges forecast-slack summaries, evaluates candidate quantities, confirms mutually preferred partners over a finite number of rounds, and reserves quantities without a central matching optimizer. Negotiation precedes current demand realization; its output is therefore a conditional commitment, not an executed transfer. After demand and due arrivals, a physical executor revalidates donor stock and receiver capacity and records actual dispatches separately from promises. This manuscript specifies the GP representation, agent–skill relationship, fixed negotiation, and simulator-grounded evaluation. The fixed gain is a forecast-based heuristic, not a calibrated value function or a guarantee of optimality. Comparative performance and industrial generalization remain empirical questions.**

***Index Terms*—Genetic programming, multi-agent coordination, inventory control, lateral transshipment, interpretable decision rules.**

I. Introduction

Multi-location inventory systems couple decisions that appear local: an order changes future availability, and a transfer improves one site's position while consuming another site's stock. Classical inventory policies provide important interpretable baselines [1], [2], while lateral transshipment introduces interactions among inventory positions, costs, and service requirements [3]. The present setting concerns sites belonging to one enterprise. They share a system-cost objective but retain local information, inventory responsibilities, and service targets. No monetary bargaining between independent firms is assumed.

GP searches executable symbolic policies rather than prescribing their final algebraic structure [4], [5]. However, directly returning one signed quantity for every site pair does not by itself define a decentralized coordination process. Simultaneous pairwise decisions may compete for the same stock; an accepted proposal may also become infeasible when demand is realized before dispatch. These issues motivate an explicit separation of policy representation, agent authority, and protocol execution.

We study three-tree GP with fixed peer negotiation. Each site executes the same evolved policy on its own observation. Replenishment, request, and offer skills use distinct GP trees, whereas routing, partner confirmation, reservations, and feasibility checks are fixed mechanisms. Offline team evaluation can be centralized without making online decisions centralized: agents exchange only specified summaries, and no agent commands another site's transfer.

The method is intended to support three investigations:

1) whether a request–offer representation with explicit commitments improves coordination relative to independent pairwise transfer rules;
2) whether GP can discover useful local replenishment, request, and offer expressions under the same fixed negotiation protocol;
3) how finite communication rounds, topology, and conditional execution affect cost and site-level service.

These are testable questions, not claims of established superiority. The implementation and editable diagrams accompany this methods draft.

II. Background and Scope

Inventory rules, simulation optimization, and symbolic policy search are complementary. Classical policies expose understandable stock targets; simulation can represent stochastic transitions and operational constraints [6]; GP can search nonlinear expressions without a differentiable simulator. A GP tree is a decision expression, not an autonomous agent. Autonomy additionally requires an observation boundary, memory, permitted messages, and authority over local actions.

Likewise, peer agreement is not synonymous with numerical consensus in networked control [7]. Here agreement means that two endpoints confirm the same conditional transfer within a finite-round protocol. We make no asymptotic convergence, strategy-proofness, privacy, or globally optimal matching claim. A skill is a callable decision module with preconditions, output bounds, and execution feedback; it is not necessarily a temporally extended reinforcement-learning option.

III. Problem and Information Structure

Consider $N \geq 2$ inventory sites on a configured undirected graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$. At the start of period t, site i observes its on-hand inventory I_i, outstanding order pipeline P_i, forecast vector $\hat{\mathbf{d}}_i$, storage capacity Cap_i, and operating cost parameters. It also maintains current-period commitments and an execution history. Let h_i, b_i, a_i, and k_i denote holding, lost-sales, unit-order, and fixed-order costs. A dispatched transfer incurs unit cost c^{tr} and fixed cost K^{tr}.

An agent cannot inspect another site's simulator object, current realized demand, future realized demand, or a full-network inventory aggregate. Neighbor communication exposes only the specified request/offer summaries, cost coefficients, and commitment updates. Such summaries are not a formal privacy mechanism.

A. Decision and Execution Timing

The timing is intentionally retained from the replenish-and-transship benchmark: agents choose orders and negotiate before observing current realized demand. The environment subsequently applies due arrivals and demand, revalidates transfers, charges costs, and advances the pipeline. Consequently, a promised quantity is not automatically deliverable.

Let $A_i(t)$ be due arrivals and $D_i(t)$ realized demand. Before physical transfers,

$$z_i(t) = \min\{Cap_i, I_i(t) + A_i(t)\} - D_i(t). \tag{1}$$

For a directed commitment $i \to j$ with planned quantity q^p, the executor can dispatch no more than

$$q^e = \min\{q^p, \max(0, z_i), \max(0, Cap_j - z_j)\}, \tag{2}$$

with z_i, z_j updated after preceding dispatches. The implementation uses a documented deterministic dispatch order. This introduces an execution-order effect when resources change after commitment; it is not claimed to remove all allocation bias. Zero, partial, and full dispatches are distinguished.

After transfers, negative inventory is charged as unmet demand and reset to zero. Positive inventory incurs holding cost. Order decisions enter the pipeline according to the benchmark lead-time convention. Receiver capacity is enforced at physical dispatch, not merely at proposal generation. The complete observable pipeline has L_o slots: its head arrives in the current period, then the pipeline shifts and the new order is appended. Thus an order chosen in period t arrives in period $t+L_o$; with $L_o = 2$, the first order arrives in period 3. This is physically equivalent to the inherited simulator's shorter queue plus separately staged action. The new policy interface explicitly exposes all outstanding orders, so baseline comparisons must align pipeline information as well as dynamics. Overflow at a capped arrival is discarded and measured.

B. Objective and Measurements

The period cost is

$$\begin{aligned} C_t = &\sum_{i \to j} \left(c^{tr} q^e_{ij} + K^{tr} \mathbb{I}[q^e_{ij} > 0]\right) \\ &+ \sum_i \left(h_i \max(z_i, 0) + b_i \max(-z_i, 0)\right) \\ &+ \sum_i \left(a_i x_i + k_i \mathbb{I}[x_i > 0]\right). \end{aligned} \tag{3}$$

The episode objective is realized cost divided by the actual number of simulated periods. Service is measured using unmet demand units and demand fill rate, not a count of zero ending inventories.

Synthetic demand uses stochastic Poisson realizations conditional on the benchmark's exposed demand means. These mean forecasts are privileged model information and should not be interpreted as evidence of real forecasting accuracy. Historical records copied into future forecast slots constitute an oracle setting; they must be explicitly marked or rejected, rather than silently presented as causal forecasting. Synthetic episodes retain the inherited reset convention: the stream seed is the supplied episode seed plus 1000. Shortened execution horizons do not regenerate a different demand table.

TABLE I
RESPONSIBILITY BOUNDARIES

Object	Responsibility
T^{ord}	Compute a raw local replenishment quantity.
T^{req}	Compute a raw support-request quantity.
T^{off}	Compute a raw support-offer quantity.
Skill wrapper	Supply allowed features, validate numeric outputs, apply nonnegativity and operational bounds.
Agent	Maintain local observation, routing identity, role, commitments, and actual execution memory.
Fixed protocol	Exchange summaries, score quantity candidates, select mutual partners, reserve and confirm.
Physical executor	Apply arrivals and demand; recheck stock/capacity; record actual transfers and costs.

IV. THREE-TREE GP WITH FIXED PEER NEGOTIATION

A. GP, Agent, and Skill Are Different Objects

An *agent* is the persistent site-level decision entity: it owns local observations, inventory responsibilities, neighbor interfaces, reservations, and execution memory. A *skill* is a bounded callable module within that agent. A *GP tree* is the symbolic expression used by a particular decision skill. GP itself is the offline search procedure that changes the expressions. None of these objects requires an LLM.

The population individual is the ordered tuple

$$\Theta = (T^{ord}, T^{req}, T^{off}), \tag{4}$$

containing exactly three trees. The same Θ is deployed at every site, but each invocation receives a different local observation. This is shared parameterization, not shared access to site states.

Table I separates evolved and fixed responsibilities.

B. Local Features and Skill Outputs

The GP feature interface contains local inventory and pipeline summaries, exposed forecasts, capacity and cost parameters, and explicitly available previous-period execution features. Features are named, so a tree's terminal meaning is independent of a legacy array position. Raw neighbor states, hidden demand samples, and global inventory totals are excluded. Site identities are used for routing and deterministic tie-breaking, not as GP decision terminals.

`PIPE` includes all previously placed outstanding orders; `DUE` is the first pipeline entry. `F` sums all exposed means, whereas `FC1` and `FC2` are the first two entries (`FC2` is zero if absent). The three `LAST` terminals refer to actual received, sent, and lost quantities in the preceding period; they start at zero. Transfer costs `TC`, `TK` are public scenario parameters. No site-ID, time-index, topology, or production-capacity terminal is used.

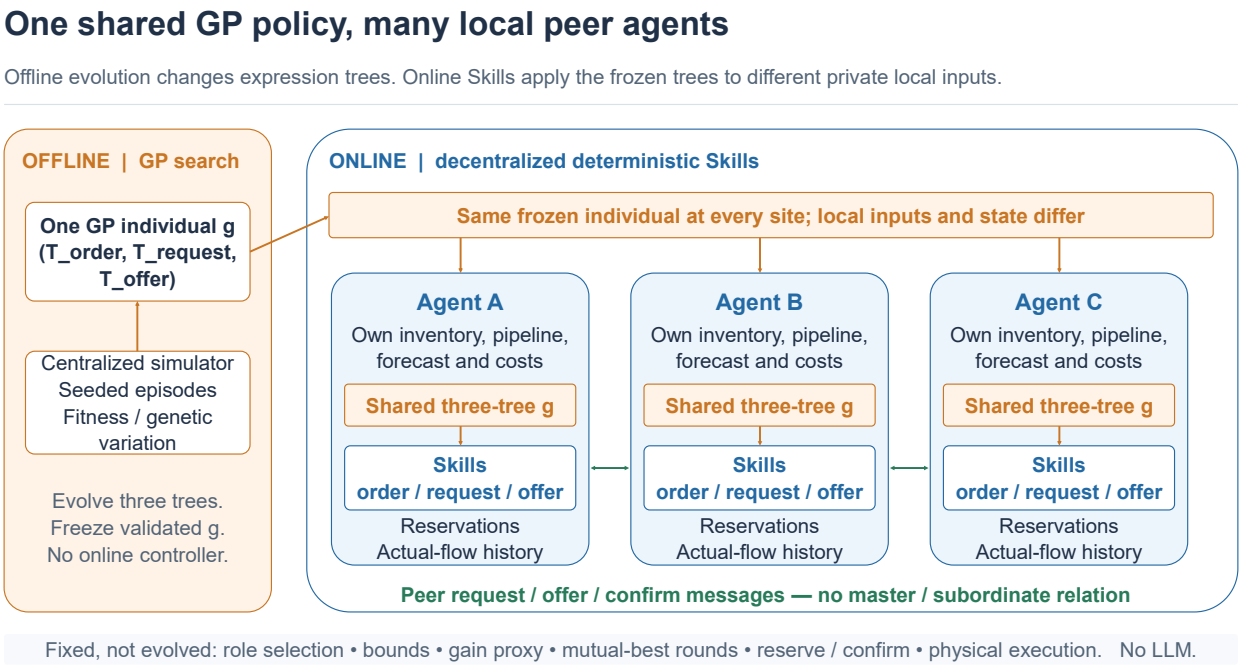


Fig. 1. Relationship between offline GP, the shared three-tree individual, site agents, and their decision skills. Fixed protocol and physical execution mechanisms are outside the evolved representation.

TABLE II
EXACT NAMED GP TERMINAL SETS

Scope	Terminals
All trees	`I, PIPE, DUE, F, FC1, FC2, CAP, H, B, RHO, LAST_IN, LAST_OUT, LAST_LOST`
Order only	`ORDER_FIXED, ORDER_UNIT`
Request only	`ROOM, DEFICIT, TC, TK`
Offer only	`AVAIL, SURPLUS, TC, TK`

Let f_i denote the allowed feature vector. Fixed wrappers convert tree outputs to bounded quantities:

$$x_i = \mathrm{clip}(T^{ord}(f_i), 0, U_i^{ord}), \tag{5}$$
$$r_i = \mathrm{clip}(T^{req}(f_i), 0, U_i^{req}), \tag{6}$$
$$s_i = \mathrm{clip}(T^{off}(f_i), 0, U_i^{off}). \tag{7}$$

The bounds are derived from the configured benchmark and current role; they are not additional learned trees. Orders are capped by storage capacity per order, not by currently empty storage space. A supply bound at negotiation time is not a guarantee of available stock after demand.

Define forecast exposure and forecast slack by

$$F_i = \sum_{\kappa} \hat{d}_i^{(\kappa)}, \qquad \rho_i = I_i + \sum P_i - F_i. \tag{8}$$

With $\epsilon = 10^{-8}$, role assignment is fixed: a site with $\rho_i > \epsilon$ can offer, one with $\rho_i < -\epsilon$ can request, and other sites are inactive. Exact bounds are

$$U_i^{ord} = Cap_i, \tag{9}$$
$$U_i^{req} = \min\{\max(0, -\rho_i), \max(0, Cap_i - I_i)\}, \tag{10}$$
$$U_i^{off} = \min\{\max(0, \rho_i), \min(Cap_i, I_i + P_i[0])\}. \tag{11}$$

Inactive-role request/offer quantities are zero. Due arrivals enter the conditional supply bound; other pipeline entries only affect forecast slack. The implementation excludes receiving and forwarding within the same period. A GP output of zero can decline participation even when the role permits it.

The observation, GP features, period role, and initial skill quantities remain frozen during negotiation. Trees are not reevaluated after each match; only message quantities and reservation-adjusted slack change. An order newly selected in the current period does not enter that period's observed pipeline or slack.

C. Expression Representation and Example

Each tree is a prefix expression with six binary primitives: `add, sub, mul, div, maximum, minimum`. Search constants are $\{0, 0.25, 0.5, 1, 2\}$. GP terminal inputs and primitive results are saturated to $[-10^6, 10^6]$. Protected division returns 1 when the denominator's absolute value is at most 10^{-8}; otherwise the quotient is saturated. Physical state, skill bounds, messages, and proxy gains are not subjected to this GP numeric saturation. Nonfinite policy values raise an explicit numerical-policy error. These are numeric expression trees, not a dimensionally typed program language.

An illustrative valid individual is

$$T^{ord}(f_i) = \max(0, F_i - I_i - \sum P_i), \tag{12}$$

$$T^{req}(f_i) = \max(0, F_i - I_i - \sum P_i), \tag{13}$$

$$T^{off}(f_i) = \max(0, I_i - F_i). \tag{14}$$

This is an expression example, not the exact initialized seed and not a reported evolved optimum. Its order and request expressions are identical but have different execution roles: the former orders material with a lead time, whereas the latter advertises demand for peer support. Fixed caps still apply. Figure 2 shows the three-tree structure.

D. Messages and Fixed Quantity Evaluation

An agent sends request or offer quantities and the forecast-slack/cost summary required by the protocol to its configured neighbors. Current-period commitments reduce the remaining advertised request/supply and adjust effective slack for subsequent negotiation rounds. Neither endpoint needs to query the other's simulator object.

The non-evolved local proxy is

$$\Phi_i(\rho) = h_i \max(\rho, 0) + b_i \max(-\rho, 0). \tag{15}$$

For a candidate quantity q from donor i to receiver j,

$$A_i(q) = \Phi_i(\rho_i - q) - \Phi_i(\rho_i), \tag{16}$$

$$B_j(q) = \Phi_j(\rho_j) - \Phi_j(\rho_j + q), \tag{17}$$

$$G_{ij}(q) = B_j(q) - A_i(q) - c^{tr} q - K^{tr} \mathbb{I}[q > 0]. \tag{18}$$

The gain is a deterministic forecast-slack heuristic. It is not an unbiased estimate of future cost, a donor-shortage probability, or an incentive compatibility condition. Its fixed definition deliberately isolates the contribution of evolving the three quantity trees.

Let $\bar{s}_i$ and $\bar{r}_j$ be remaining quantities after reservations. The protocol evaluates a small grid:

$$\mathcal{Q}_{ij} = \left\{ \eta \min(\bar{s}_i, \bar{r}_j) : \eta \in \{0, \tfrac{1}{4}, \tfrac{1}{2}, \tfrac{3}{4}, 1\} \right\}. \tag{19}$$

For each ascending grid fraction, the implementation computes $q = \min\{q_{\max}, \mathrm{round}(\eta q_{\max}, 6)\}$, where $q_{\max} = \min(\bar{s}_i, \bar{r}_j)$. Rounding uses ties-to-even; clipping after rounding cannot exceed the remaining budget. Candidates are deduplicated within 10^{-8}. Orders, initial skill budgets, and actual dispatch quantities are otherwise not rounded. Only $q > 10^{-8}$ and $G_{ij}(q) > 10^{-8}$ are eligible. Common ranking is gain descending, donor ID ascending, receiver ID ascending, then quantity descending; gains are compared exactly.

The conservative surplus/deficit caps have an important consequence: no candidate crosses the proxy's donor or receiver breakpoint. Thus,

$$A_i(q) = -h_i q, \qquad B_j(q) = b_j q, \tag{20}$$

$$G_{ij}(q) = (h_i + b_j - c^{tr})q - K^{tr}. \tag{21}$$

Whenever an eligible candidate exists, the largest available grid quantity is preferred. The grid is therefore not discovering an interior nonlinear risk optimum in this version. GP still controls advertised budgets, which change edge limits, gains, partner choices, and commitments.

Algorithm 1 Fixed peer negotiation in one period

```
Each agent computes local x_i, r_i, s_i using Θ.
Assign roles and initialize empty reservation ledgers.
for r = 1, ..., R do
    Exchange remaining quantities and summaries on E.
    Each endpoint evaluates incident quantity grids using (18).
    Each agent selects its highest-ranked positive-gain edge.
    Confirm only mutually selected edges; reserve agreed quantities.
    Update remaining quantities and effective forecast slack.
    if no edge is confirmed then
        Stop negotiation.
    end if
end for
Return orders and conditional commitments.
```

E. Finite-Round Mutual Matching and Reservations

Independent bilateral acceptance can over-commit a donor. Instead, in each round every active agent chooses its highest-ranked eligible incident edge. An edge is confirmed only when its endpoints choose one another. A common ranking combines predicted gain and deterministic identifiers, so endpoints agree on tie resolution.

Each site can confirm at most one partner in a round, but may confirm different partners over successive rounds while resources remain. Confirmation immediately updates both ledgers and the advertised remaining quantities. The protocol terminates after the configured round limit or when no additional mutual match exists. A simulator can deliver messages synchronously without acting as a global matching optimizer.

Ledger invariants bound total planned outgoing supply and planned incoming requests by each site's advertised remaining budgets. They prevent duplicate commitments; they do not guarantee that post-demand inventory can fulfill every promise. No global optimality or network-consensus convergence theorem is asserted.

F. Reservation Invariants and Local Authority

Write $\mathcal{C}$ for confirmed commitments in the current period. For a donor and receiver respectively, remaining budgets are

$$\bar{s}_i = s_i - \sum_{(i,j,q) \in \mathcal{C}} q, \tag{22}$$

$$\bar{r}_j = r_j - \sum_{(i,j,q) \in \mathcal{C}} q. \tag{23}$$

A new confirmation uses $0 < q \le \min(\bar{s}_i, \bar{r}_j)$ and updates both endpoints together. Starting from empty ledgers, induction over confirmations gives

$$\sum_{(i,j,q) \in \mathcal{C}} q \le s_i, \qquad \sum_{(i,j,q) \in \mathcal{C}} q \le r_j. \tag{24}$$

The first inequality concerns each donor separately and the second each receiver separately. They are promise-budget invariants, not assertions about inventory after future demand.

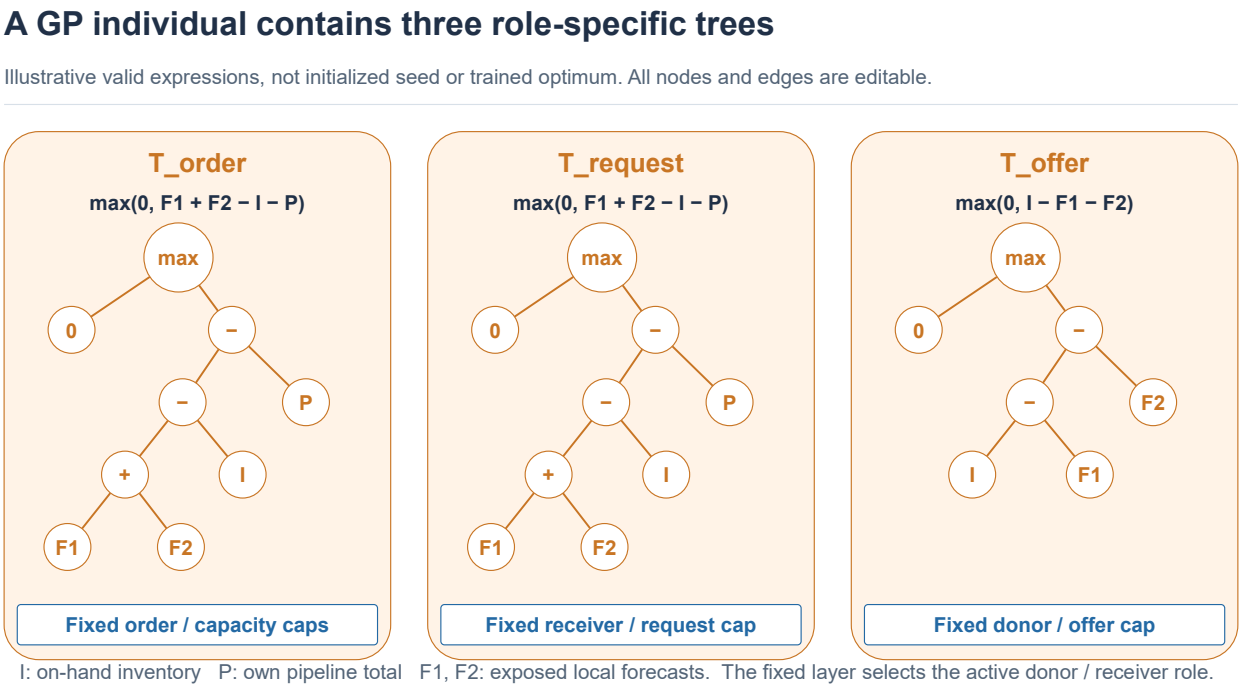


Fig. 2. An illustrative three-tree individual. Trees encode expressions; skills supply semantics, bounds, and lifecycle. The diagrams remain editable in the accompanying PowerPoint deck.

Effective slack is updated consistently with commitments: donating q subtracts q from the donor's negotiation slack, while receiving q adds it to the receiver's negotiation slack. An offer/request message does not itself reserve stock. A partner preference does not itself confirm a transfer. Only a mutually accepted confirmation changes the reservation ledger; only physical dispatch changes actual transfer memory. This lifecycle prevents an agent from treating an unaccepted proposal as already received material.

G. Conditional Dispatch and Execution Memory

After current demand and due arrivals become available to the environment, the fixed executor processes commitments by negotiation round, donor ID, and receiver ID, using (2). Negotiation is not retrospectively rerun with realized demand. For every commitment, it records planned and dispatched quantities, full/partial/cancelled status, and the reason for any reduction. Only dispatched quantities update the actual transfer history. Costs are charged on actual dispatches; an approved but cancelled promise does not incur a fictitious physical transfer cost.

H. Offline GP Search and Fitness

Initialization creates valid three-tree individuals, including interpretable seed rules. The supplied seed consists of `maximum(0.0, sub(F, add(I, PIPE)))`, `DEFICIT`, and `SURPLUS`, in order/request/offer order. Tournament selection and elitism operate on complete individuals. Crossover exchanges subtrees only between matching roles; mutation changes one selected role's expression. Height/node limits control bloat. The fixed negotiation and execution machinery is never crossed over or mutated.

All competing candidates and retained elites use the same fixed training scenario/seed batch across generations. This version does not rotate batches. Unique generation champions form the candidates for selection on a disjoint validation batch. The chosen policy is saved before held-out testing. Let D_{ei} and L_{ei} be realized demand and final unmet demand for site i in episode e. Define

$$\widehat{FR}_{ei} = \begin{cases} 1 - L_{ei}/D_{ei}, & D_{ei} > 0, \\ 1, & D_{ei} = 0. \end{cases} \tag{25}$$

With a common configured site service target τ, exact batch fitness is

$$\mathcal{F}(\Theta) = \left(V(\Theta), \overline{J}(\Theta), |\Theta|\right), \tag{26}$$

$$V = \max_{e} \sum_{i} \max(0, \tau - \widehat{FR}_{ei}), \tag{27}$$

$$\overline{J} = \frac{1}{M} \sum_{e=1}^{M} \frac{C_e}{T_e}. \tag{28}$$

There is no division of cost by site count. Smaller lexicographic values are preferred; exact ties use canonical tree-string tuples. The service constraint is site-wise, so a neglected site cannot be hidden by a high network-aggregated fill rate. This remains empirical constraint handling, not a statistical safety certificate.

Variation selects one homologous module. Structurally excessive children revert to their corresponding parent. Search

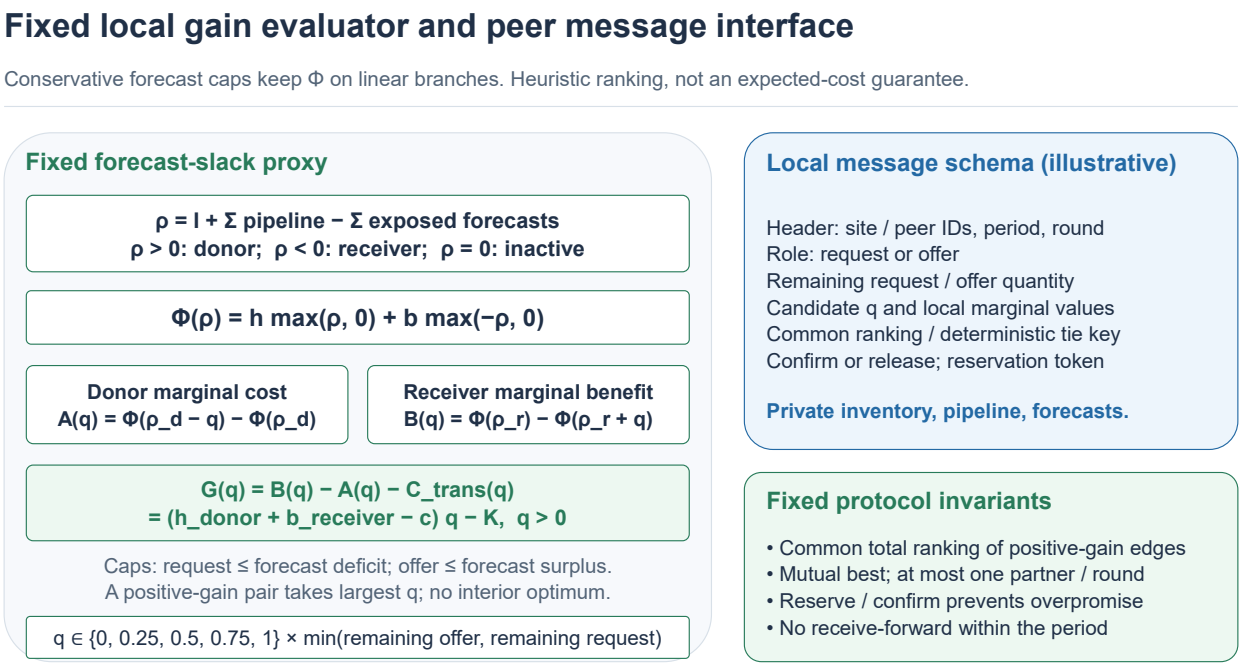


Fig. 3. The fixed evaluator uses advertised summaries to score a finite quantity grid. It is not an evolved fourth tree or an online LLM.

TABLE III
IMPLEMENTED SEARCH DEFAULTS

Setting	Default
Population / replacement generations	20 / 5
Tournament size / elites	3 / 1
Crossover / mutation / reproduction	0.80 / 0.15 / 0.05
Initial half-and-half depth	0–2
Mutation grow depth	0–2
Maximum height / nodes per tree	5 / 63
Negotiation rounds / site fill target	3 / 0.95
Training seeds	11, 12
Validation seeds	1001, 1002
Held-out seeds	2001, 2002

is serial; the smoke command uses population six and one replacement generation without shortening the default 64-period scenario. These are prototype defaults, not tuned experimental budgets.

I. What Is Deliberately Not Learned

This version does not evolve message syntax, communication topology, quantity-grid definition, gain equations, matching protocol, or physical constraints. It does not use five-tree value learning, partner-pool coevolution, or counterfactual credit assignment. Those are separate extensions requiring new ablations. Keeping these boundaries explicit makes the first three-tree prototype reproducible and its empirical effects attributable.

J. Computational and Communication Structure

Let L be the total number of nodes in the three expressions and K_q the number of candidate quantities per edge. Local tree execution is linear in the visited expression size. Across all sites and negotiation rounds, a period requires on the order of

$$O(NL + R|\mathcal{E}|K_q) \tag{29}$$

arithmetic work, apart from state updates and physical execution. Summary exchange, preferences, and confirmations require $O(R|\mathcal{E}|)$ bounded-size messages in a synchronous implementation. These are structural counts, not measured

Algorithm 2 Three-tree GP policy search

1: Initialize population of $(T^{ord}, T^{req}, T^{off})$.
2: Evaluate the initial population on a common training batch.
3: **for** each generation **do**
4: Rank the evaluated population by violation, cost, and size.
5: Retain elites and select parents for offspring.
6: Apply same-role crossover and module mutation.
7: Enforce structural and numerical validity constraints.
8: Evaluate offspring on the same current training batch.
9: Form the next population from evaluated offspring and elites.
10: **end for**
11: Select among unique generation champions on validation trajectories.
12: Freeze and serialize its three expressions and configuration.
13: Evaluate once on the held-out test protocol; deploy local skills.

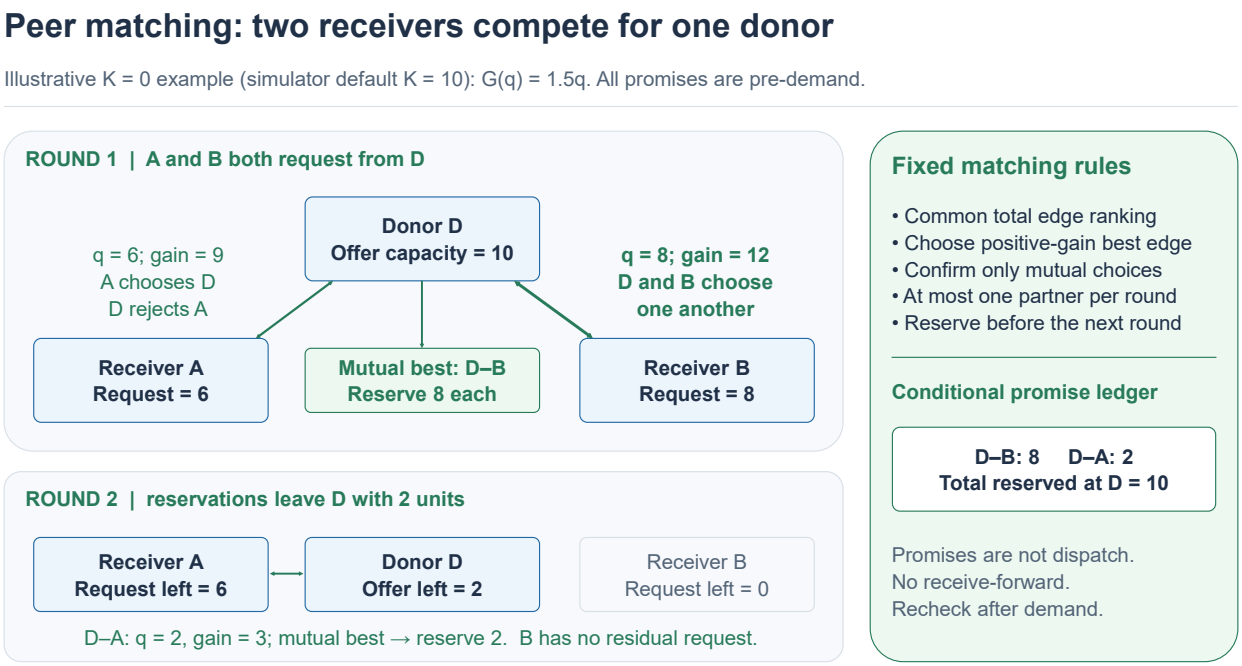


Fig. 4. Peer negotiation with competing support requests. Local partner choices and reservations replace a central matching decision. Multiple rounds allow a donor to allocate residual supply.

latency guarantees. For a complete graph, $|\mathcal{E}| = N(N-1)/2$; parameter sharing alone does not remove quadratic all-to-all communication. Sparse-topology and round-budget experiments are therefore essential to any scalability claim.

V. Evaluation Protocol

This manuscript supplies an executable methods prototype, not a completed comparative study. Training uses common scenario/seed batches across competing individuals. Validation selects among candidate policies; final test trajectories do not participate in selection. Demand repetitions measure a frozen policy's variability and must not be counted as independent GP search runs.

Required baselines include tuned local inventory rules, no lateral transshipment, the copied two-tree replenish/transship GP, and the same three-tree representation with independent pairwise acceptance instead of mutual matching. A fixed hand-written request/offer policy isolates the benefit of GP search from the benefit of the protocol. Information, timing, action bounds/precision, scenario distribution, and simulation budgets must be aligned. The new skill wrappers do not apply the legacy logistic action map or its two-decimal rounding.

Ablations should vary negotiation rounds, graph connectivity, request/offer evolution, and reservation handling. An unconstrained oracle-information controller may be used as a reference, but is not automatically an optimal bound. Stress tests should include multiple receivers competing for one donor, asymmetric costs, demand shocks, and larger networks.

Report average cost and its components, per-site fill rate, unmet units, planned and executed transfer quantities, full/partial/cancelled commitment counts, and protocol communication. Report search wall time, episode evaluations, tree sizes, independent search seeds, and paired test uncertainty. Small smoke runs validate execution and reproducibility only; they cannot establish an algorithmic advantage.

VI. Limitations and Conclusions

Three-tree GP replaces a general code generator with a restricted symbolic search space. The fixed protocol intentionally remains outside evolution, which improves attribution but limits what coordination patterns can be discovered. The gain proxy ignores much of the multi-period stochastic problem; positive predicted gain is not a guarantee of positive realized benefit. Shared rules do not establish performance with arbitrary heterogeneous partners. Finite mutual matching need not solve the global allocation problem, and conditional commitments can fail after demand.

The resulting research object is nevertheless precise: local inventory agents use GP-parameterized skills to produce requests and offers; a fixed peer protocol converts compatible choices into conditional commitments; a separate executor enforces physical limits and records realized outcomes. The next empirical task is to test whether this separation yields useful cost–service tradeoffs under equal information and evaluation budgets.

References

[1] P. H. Zipkin, *Foundations of Inventory Management*. McGraw-Hill, 2000.
[2] S. Axsater, *Inventory Control*, 3rd ed. Springer, 2015.

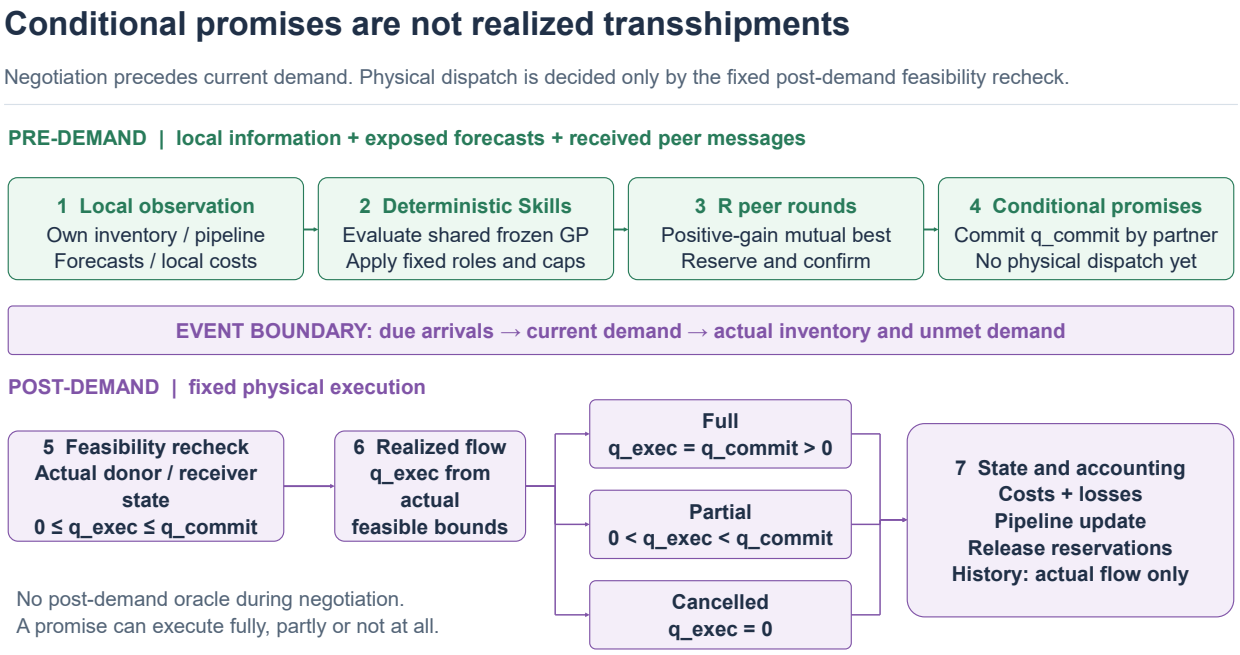


Fig. 5. Pre-demand commitments and post-demand execution are separate events. Agents do not use future realized demand when negotiating.

[3] C. Paterson, G. Kiesmuller, R. Teunter, and K. Glazebrook, "Inventory models with lateral transshipments: A review," *European Journal of Operational Research*, vol. 210, no. 2, pp. 125–136, 2011.

[4] J. R. Koza, *Genetic Programming: On the Programming of Computers by Means of Natural Selection*. MIT Press, 1992.

[5] R. Poli, W. B. Langdon, and N. F. McPhee, *A Field Guide to Genetic Programming*. Lulu, 2008.

[6] S. Amaran, N. V. Sahinidis, B. Sharda, and S. J. Bury, "Simulation optimization: A review of algorithms and applications," *Annals of Operations Research*, vol. 240, no. 1, pp. 351–380, 2016.

[7] R. Olfati-Saber, J. A. Fax, and R. M. Murray, "Consensus and cooperation in networked multi-agent systems," *Proceedings of the IEEE*, vol. 95, no. 1, pp. 215–233, 2007.

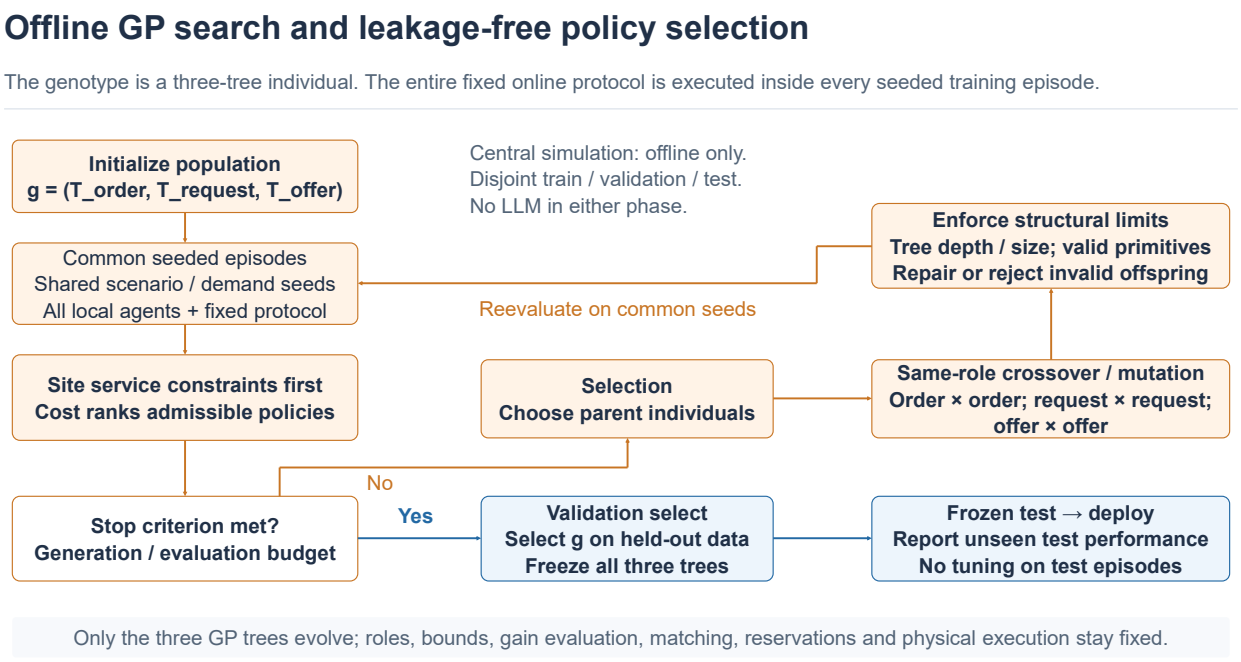


Fig. 6. Offline GP search and separate validation/test evaluation. Deployment executes frozen expressions; it performs neither evolution nor LLM inference.